

ECE421: Introduction to Machine Learning
Programming Assignment 3
Assigned: March 2, 2024; Due: March 18, 2024 @ 11:59 p.m.

Objectives

The purpose of this assignment is to investigate the classification performance of neural networks. You will implement, train and evaluate your neural network models using Pytorch on the notMNIST dataset. The functions are provided and **you will need to fill-in after the `TODO` comment instead of coding them from scratch**. In particular, in the first part, you will implement a fully connected neural networks (FNN). In the second part, you will implement the model training loop. You will also be asked to answer several questions related to your implementations. You are encouraged to look up Pytorch documentation for useful utility functions, at: <https://pytorch.org/docs/stable/index.html>. You can also refer to the demo covered in the lecture. The notMNIST dataset is stored as `notMNIST.npz` inside in the assignment folder. We highly recommend students to look at the tutorial files for this assignment.

To avoid any potential installation issue, you are encouraged to develop your solution using Google Colab notebooks. It is highly recommended that you train the neural network using a GPU.

Unlike in the tutorial video which implemented convolutional neural networks, we simply implemented a fully-connected network for the image recognition task.

Requirements

In your implementations, please use the function prototype provided (i.e. name of the function, inputs and outputs) in the detailed instructions presented in the remainder of this document. We will be testing your code using a test function that will evoke the provided function prototype. If our testing file is unable to recognize the function prototype you have implemented, you can lose significant portion of your marks. In the assignment folder, the following files are included in the `starter_code` folder:

- `NeuralNetPyTorch.py`
- `notMNIST.npz`: this is a dataset and you can upload this file to Google Colab. We provide functions to read the data in `NeuralNetPyTorch.py`.

These files contain the test function and an outline of the functions that you will be implementing. You also need to submit a separate `PA3_qa.pdf` file that answer questions related to your implementations.

Abbreviations

Following the definition in our code, we use the following abbreviations in this document:

- `BATCH_SIZE` refers to the batch size of images we are using for training. In this assignment, we set `BATCH_SIZE=32`.
- `F` is an abbreviated import name of `torch.nn.functional`.
- `nn` is an abbreviated import name of `torch.nn`.

Preliminaries

In this part, we explain several helper functions/classes in the `NeuralNetPyTorch.py` file. **Do not modify any of these functions/classes.**

- Function 1: `loadData(datafile="notMNIST.npz")`
 - Inputs: `datafile`
The default input is `"notMNIST.npz"`, which is the dataset we are using in this assignment.
 - Output: `trainData`, `validData`, `testData`, `trainTarget`, `validTarget`, `testTarget`
The outputs are images and annotations in the form of Numpy matrices. `trainData` and `trainTarget` are the images and annotations for training. Similarly, `validData` and `validTarget` are the images and annotations for validation and `testData` and `testTarget` are the images and annotations for testing.
 - Functionality:
This function loads the notMNIST dataset and splits it into training, validation and testing set.
- Class 1: `class notMNIST(Dataset)`
 - Functionality:
This class implements a PyTorch Dataset class for the notMNIST dataset.

Read the description of the function below as you will use it in this assignment

- Function 2: `def experiment(learning_rate=0.0001, num_epochs=50, verbose=False):`
 - Inputs: `learning_rate`, `num_epochs`, `verbose`
The input `learning_rate` is a scalar of type `float` that specifies the learning rate. The input `num_epochs` is an integer that specifies how many times your model will train through the whole training set. The input `verbose` is a Boolean that will let you print the training process if is `True` and nothing if is `False`.
 - Output: trained model and training history
This function returns a trained model, which is a `torch.nn.Module` class and its training history. The training history is a dictionary that contains the accuracy on training, validation and test sets at each epoch, which is returned after calling the `train` function. **Note: You will implement the `train` function in Part 3.**
 - Functionality:
This function will build your model, train it on the notMNIST dataset using the specified hyperparameters and return a trained model with its training history.

Part 1: Fully Connected Neural Networks

In this part, you will be implementing a Fully Connected Neural Network using PyTorch. The model is defined in `class FNN(nn.Module)`. The neural network architecture that you will be implementing is as the following order:

1. An input layer for the images of size $(\text{BATCH_SIZE} \times 1 \times 28 \times 28)$. The second dimension is the image channel, which is 1 in this case since we are using gray-scale images. The third and forth dimension are the width and height of the input respectively. We will transform the $(\text{BATCH_SIZE} \times 1 \times 28 \times 28)$ matrix into a batch of 1D arrays of size $(\text{BATCH_SIZE} \times 784)$.
2. A fully connected layer followed by a ReLU activation. The size of the weight matrix is 784×10 where 784 is the size of the input array and 10 is the size of the 1st hidden layer.
3. A fully connected layer followed by a ReLU activation. The size of the weight matrix is 10×10 where 10 is the size of the 1st hidden layer and 10 is the size of the 2nd hidden layer.

4. A fully connected layer (**without softmax activation**). The size of the weight matrix is 10x10 where 10 is the size of the 2nd hidden layer and 10 is the size of the output layer.

Specifically, you will be implementing two functions in `class FNN(nn.Module)`, which are detailed in the following:

- Function 1: `def __init__(self)`
 - Inputs: This function does not require any input
 - Output: This function does not return any output
The purpose of this function is to setup the variables for this class. As shown in the tutorial, you need to define all the layers you will be using in this function.
 - Function implementation considerations:
You will use the following PyTorch functions to set up the layers in your network: `nn.Linear()`. The `nn.Linear()` function is a fully connected layer and is defined by the dimension of the input and output. For example, `nn.Linear(3, 4)` is a fully-connected layer for an input of dimension `(BATCH_SIZE, 3)` and an output of dimension `(BATCH_SIZE, 4)`.
- Function 2: `forward(self, x)`
 - Inputs: `self, x`
The input `x` is the batch of images of size `(BATCH_SIZE, 1, 28, 28)`. The input `self` represents the instance of the class.
 - Output: `out`
This function computes the logits for each image in the batch. The output has a size of `(BATCH_SIZE, 10)`, where each `(i, j)` position represent the class-logit score `j` of image `i`. The order of the layers (and activations) in this function must follow the described network architecture above.
 - Function implementation considerations:
You will find the following PyTorch functions helpful: `torch.flatten(x, start_dim=1)`, `F.relu()`. The `torch.flatten(x, start_dim=1)` operation will flatten your input `x` of size `(BATCH_SIZE, 1, 28, 28)` to the size of `(BATCH_SIZE, 784)`. Functions `F.relu()` applies the `ReLU()` activation to the input respectively.

The following is the mark breakdown for Part 1:

- Test file successfully runs implemented function: 15 marks
- Output is close to the expected output from the test file: 15 marks

Part 2: Model Training and Experiments

In this part, you will be implementing the training procedure for your neural networks. You will use the cross-entropy loss, Adam optimization method. Implement (Complete) the following functions:

- Function 1: `def get_accuracy(model, dataloader):`
 - Inputs: `model, dataloader`
`model` is an instance of the neural network class. `dataloader` is an instance of the `notMNIST` dataloader (see the function `experiment`).
 - Output:
The output of this function is a scalar, which is the accuracy over all images in `dataloader`.
 - Function implementation considerations:
This function will calculate the classification accuracy over the `dataloader` we specified. It will be called whenever your model finishes one training epoch in the `train` function. Remember to set `model.eval()` to get the model into evaluation mode.

- **Function 2:** `def train(model, device, learning_rate, train_loader, val_loader, test_loader, num_epochs=50, verbose=False):`
 - **Inputs:** `model`, `device`, `learning_rate`, `train_loader`, `val_loader`, `test_loader`, `num_epochs=50`, `verbose`
`model` is an instance of the neural network class. To train the neural networks with GPU, `device` should be `"cuda:0"` as shown in the `experiment` function. The `learning_rate` is a float scalar specifying the learning rate at training. `train_loader`, `val_loader`, `test_loader` are the `notMNIST` dataloader for training, validation and testing set respectively. `num_epochs` is the number of passes through the dataset we would like in our training (default is 50 in this assignment). Finally, set `verbose` to `True` will print out the training progress.
 - **Output:** `model`, `acc_hist`
This function will return a trained `model` and the evolution of training, validation and testing accuracy.
 - **Function implementation considerations:**
First, you need to specify the cross-entropy loss through the `criterion` variable that we will use as an optimization objective. After that, set the AdamW optimizer through the `optimizer` variable (PyTorch Link) using the learning rate and the model's weights (to be optimized). After that, complete the backpropagation process (forward, backward, update weights) in the training loop in the function. Remember to set the weight's gradients to zero.

Complete the following experiments, write and save your answer in a separate `PA3-qa.pdf` file. You will also need to write functions that prints/plots these experiment in your `NeuralNetPyTorch.py` submission. Remember to submit this file together with your code. Use the `experiment` function to complete these experiments. **For simplicity, we will discard the validation accuracy since fine-tuning the hyper-parameters for neural networks is a very time-consuming process.** To run the experiment, use the provided `experiment` function.

- **Experiment:** In this experiment, you will train the FNN in the image recognition task. Train your FNN model for a batch size of 32 (default) for 50 epochs (default) and the AdamW optimizer (Link to AdamW) for learning rate of 0.1, 0.01, 0.001, 0.0001. Plot and compare the training/testing history of these models, namely we want to see how the training progresses with different learning rates. The function name for this experiment will be `compare_lr()`. Comment on what you observed from the experiment and explain.

The following is the mark breakdown for Part 2:

- Test file successfully runs implemented function: 6 marks
- Output is close to the expected output from the test file: 18 marks
- Questions are answered correctly: 6 marks